



InsightLoop

Product Requirements Document

Project health reviews powered by structured human input and governed AI



SELISE Group AG
The Circle 37
8058 Zürich-Flughafen
Switzerland
CH +41 44 805 8044
www.selisegroup.com

Document Information

Document	Value
Status	Finalized requirements baseline
Version	1.0
Date	18 August 2026
Audience	Product, Engineering, Design, QA, Operations, Department Leadership
Scope	Internal company platform; project-level reporting only

Decision baseline This PRD preserves the finalized workflow. Items not explicitly fixed—health-score thresholds/weights, reminder timing, and operational time limits—are configurable policies and must not weaken mandatory coverage, privacy, or approval controls.

1. Executive Summary

InsightLoop is an internal, multi-department platform for recurring and ad-hoc project health reviews. Departments configure reusable evaluation and report templates, select internal participants with flexible project-role context, and launch a single review template across multiple functions. AI creates role-relevant questionnaires, conducts bounded follow-up interviews, compares participant perspectives, calculates explainable project health, and drafts one consolidated report. A human Project Owner edits the draft and a Unit Admin must approve it before internal distribution in-app, by PDF, or email.

Core product promise Turn independent, role-specific project knowledge into a governed, evidence-linked, approved project health report without maintaining separate questionnaires for every role.

2. Problem and Opportunity

Software delivery organizations manage many projects across departments and accounts. Monthly and quarterly reviews are costly to coordinate, inconsistent across projects, and vulnerable to single-person bias. Static questionnaires either become generic or proliferate by role. Manual synthesis often hides disagreement and weakens traceability.

- Standardize the review discipline while allowing department and project-specific criteria.
- Capture independent perspectives from Dev, QA, BA, PM, Tech Lead, UX, and other functions.
- Use AI as both a bounded interviewer and a cross-participant analyst.
- Preserve human accountability through editing, mandatory approval, versioning, and auditability.

3. Goals and Non-Goals

3.1 Goals

- Support Company → Department → Account → Project → Review Instance hierarchy with strict department isolation.
- Achieve 100% assignment coverage of mandatory template questions before launch.
- Generate adaptive questionnaires from one versioned template using participant role context.
- Produce explainable, evidence-linked cross-participant analysis and project health scoring.
- Require Unit Admin approval before any report becomes final or is distributed.

- Maintain reproducible project review history using immutable template/report-template versions.

3.2 Non-Goals

- External/customer questionnaire participation or report recipients in MVP.
- Account-, department-, company-, or portfolio-level aggregate reporting; reporting is project-level only.
- Automatic publication of AI output without human approval.
- Unlimited conversational interviewing or general-purpose AI chat.
- Automated performance appraisal of individuals.
- Period-over-period trend analytics and broad third-party integrations in MVP.

4. Personas and Roles

Role / Concept	Primary responsibility	Scope
Admin	Company governance, departments, users, company standards, security and oversight.	Company-wide
Unit Admin	Administers own department, templates, projects, reviews, and mandatory final-report approval.	Own department
User	Internal employee who may participate in reviews or receive final reports.	Authorized/selected projects
Project Owner	User assigned at project level to operate reviews and edit assigned project information.	Assigned project(s)
Participant	Selected internal user who independently answers a questionnaire with one or more functions.	Own questionnaire
Recipient	Selected internal user authorized to receive the approved final report.	Specific final report

Project Owner is an assignment layered on the User role, not a global base role. A person may have multiple participant functions (for example BA + QA or Developer + Tech Lead), and these functions may vary by review.

5. Core Concepts and Data Hierarchy

Company → Department → Account → Project → Review Schedule → Review Instance → Questionnaire / Response → Analysis → Report Revision → Distribution Event

Concept	Definition / relationship
Department boundary	Security tenant for accounts, projects, templates, reviews, responses, reports, and recipients.
Review Template	Versioned criteria and questions used for coverage and questionnaire generation.
Report Template	Separately versioned structure used to generate the consolidated report.

Concept	Definition / relationship
Review Schedule	Monthly, quarterly, custom recurring, or ad-hoc configuration; recurring schedules require an end date.
Review Instance	A project review for a defined period, frozen to specific template versions.
Participant Context	One or more project functions recorded when a user is selected.
Coverage Map	Assignment of every mandatory criterion/question to at least one participant.
Report Revision	Draft, reviewed, approved, and immutable final representations with provenance.

6. End-to-End Workflows

6.1 Template authoring and publishing

1. Admin or Unit Admin creates a template manually, duplicates an existing version, or uploads a PDF/DOCX for AI-assisted extraction.
2. System produces an editable draft with criteria, questions, types, guidance, and suggested role relevance.
3. Author reviews and corrects the draft, then publishes an immutable version.
4. Subsequent edits create a new draft/version; existing reviews remain bound to their original version.
5. The same versioning lifecycle applies separately to report templates.

6.2 Review setup and launch

6. Project Owner or Unit Admin creates an ad-hoc review or a recurring schedule with start date, frequency, and required end date.
7. User selects review-template and report-template versions, internal participants, each participant's functions, and internal final-report recipients.
8. AI proposes role-relevant assignments and an optional preview.
9. System calculates coverage. Launch remains blocked until every mandatory template criterion/question has at least one participant (100%).
10. On launch, the system creates private participant questionnaires and sends notifications.

6.3 Questionnaire and adaptive interview

11. Participant opens only their own questionnaire and answers independently.
12. AI may ask a follow-up only for material risk, ambiguity, contradiction, missing rationale, or insufficient evidence.
13. AI stops when sufficient information exists or at the first applicable cap: 3 follow-ups per template question or 10 follow-ups per participant questionnaire.
14. Participant saves/resumes and submits. Authorized reopening is audited and returns the review to incomplete collection.

6.4 Analysis, approval, and distribution

15. After all required questionnaires are complete, AI compares responses for agreement, disagreement, complementary information, blind spots, risk, and inconsistency.
16. System calculates explainable health scores and drafts one consolidated project report using the bound report-template version.

17. Project Owner and Unit Admin review/edit the draft with change history and evidence traceability.
18. Project Owner submits for approval. Unit Admin approves or returns it with comments.
19. Only after approval does the system lock a final revision and distribute it to selected internal recipients in-app, as PDF, and/or by email.
20. The closed instance remains in project history.

7. Detailed Functional Requirements

7.1 Organization, identity, and isolation

- FR-ORG-01 Support Admin, Unit Admin, and User base roles.
- FR-ORG-02 Enforce department isolation on every object and query; Admin is the only company-wide role.
- FR-ORG-03 Allow Admin and Unit Admin to create accounts and projects within their scope.
- FR-ORG-04 Allow Admin/Unit Admin to assign Project Owners; allow Project Owners to edit assigned project information.
- FR-ORG-05 Restrict the platform, participants, and recipients to active internal users.
- FR-ORG-06 Audit authentication-sensitive, configuration, content, approval, and distribution events.

7.2 Template and report-template management

- FR-TPL-01 Provide reusable basic templates and department-owned templates.
- FR-TPL-02 Support criteria, evaluation areas, questions, response types, required flags, guidance, scoring configuration, and role relevance.
- FR-TPL-03 Publish immutable versions; reviews retain exact bound versions.
- FR-TPL-04 Generate an editable draft template from PDF/DOCX; never auto-publish extracted content.
- FR-TPL-05 Manage and version report templates separately from questionnaire templates.

7.3 Review scheduling and setup

- FR-REV-01 Support monthly, quarterly, custom recurring, and ad-hoc project reviews.
- FR-REV-02 Require an end date for every recurring schedule.
- FR-REV-03 Generate distinct future review instances and stop at the end date.
- FR-REV-04 Allow flexible participant combinations and multiple functions per participant.
- FR-REV-05 Maintain separate participant and final-recipient lists.
- FR-REV-06 Bind each review instance to a project and prohibit non-project report scope.

7.4 Coverage and questionnaire generation

- FR-COV-01 Map all mandatory template criteria/questions to selected participants.
- FR-COV-02 Block launch unless coverage equals 100%; users cannot accept or waive a gap.
- FR-COV-03 Explain gaps and recommend participant/function changes.
- FR-COV-04 Offer, but do not require, a pre-launch assignment preview or manual adjustment.
- FR-QST-01 Generate a private role-adaptive questionnaire from one underlying template.
- FR-QST-02 Preserve question provenance to the template version and criterion.
- FR-QST-03 Support save/resume, submission, due dates, reminders, and audited reopen.

7.5 AI analysis and scoring

- FR-AI-01 Conduct bounded contextual follow-ups under the guardrails in Section 8.
- FR-AI-02 Wait for all required submissions before official consolidated analysis.

- FR-AI-03 Detect agreement, disagreement, complementary evidence, blind spots, risks, and inconsistencies.
- FR-AI-04 Link findings and scores to evidence accessible to authorized reviewers.
- FR-AI-05 Calculate health scores using versioned weights/thresholds and show rationale, confidence, and limitations.

7.6 Report review, approval, and distribution

- FR-RPT-01 Generate one consolidated draft project report.
- FR-RPT-02 Permit authorized human editing with provenance and revision history.
- FR-RPT-03 Require Unit Admin approval; Project Owner cannot bypass or self-finalize.
- FR-RPT-04 Support return-for-revision with comments and resubmission.
- FR-RPT-05 Lock approved final revisions; later corrections create a new revision.
- FR-DST-01 Distribute final reports in-app, by PDF, and by email to selected internal recipients only.
- FR-DST-02 Log recipient, channel, timestamp, status, and retry outcome.

7.7 History, status, and notifications

- FR-HIS-01 Maintain project-level review history and report revisions.
- FR-HIS-02 Provide status dashboards and filters within authorization scope.
- FR-HIS-03 Show completion status during collection without revealing raw peer answers.
- FR-HIS-04 Notify relevant users at launch, reminder, completion, approval request, return, approval, finalization, and distribution.

8. AI Behavior and Guardrails

Area	Requirement
Purpose limitation	AI operates only within the selected template, participant context, review evidence, and report template.
Follow-up trigger	Ask only for material risk, ambiguity, contradiction, missing rationale, or insufficient evidence.
Bounded depth	Maximum 3 AI follow-ups per template question and 10 per participant questionnaire, whichever is reached first.
Coverage	AI follow-ups supplement but never replace mandatory template-question coverage.
Privacy	Do not reveal or quote another participant's answers during collection.
Evidence	Findings and scores must be traceable to source answers; uncertainty and missing evidence must be stated.
No fabrication	Do not invent facts, participants, evidence, commitments, owners, dates, or scores unsupported by configured rules.
Human control	Label AI-generated content; allow authorized editing; never finalize or distribute without Unit Admin approval.
Safety and tone	Use neutral, non-accusatory language; report conflicting perspectives without assigning personal blame.
Version reproducibility	Record model/prompt/configuration identifiers needed for audit and controlled regeneration.

9. Status Model

Status	Entry condition	Permitted transition
Draft	Review setup started.	Coverage Blocked / Ready to Launch / Cancelled
Coverage Blocked	Coverage < 100%.	Draft / Ready to Launch
Ready to Launch	Coverage = 100% and setup valid.	Collecting
Collecting	Questionnaires sent.	Collection Complete / Cancelled
Collection Complete	All required submissions complete.	Analyzing
Analyzing	AI analysis running.	Report Draft
Report Draft	Draft generated or revision reopened.	Pending Approval
Pending Approval	Submitted to Unit Admin.	Revision Required / Approved
Revision Required	Unit Admin returned with comments.	Pending Approval
Approved	Unit Admin approval recorded.	Final
Final	Immutable final revision created.	Distributed
Distributed	At least one configured delivery completed/logged.	Closed
Closed	Workflow complete.	—
Cancelled	Stopped with reason.	—

10. Permissions Model

Permissions combine base role, department boundary, project assignment, and review-specific selection. The detailed matrix is supplied in the companion workbook. The controlling rules are:

- Admin: company-wide governance and oversight.
- Unit Admin: full departmental administration and the mandatory report approver.
- User: no blanket project administration; access derives from participation, recipient authorization, or Project Owner assignment.
- Project Owner: may edit assigned project information and operate its reviews, but cannot assign their own scope or bypass Unit Admin approval.
- Participants see only their own questionnaire and do not see peer responses during collection.
- Recipients are internal only and receive only approved final reports for which they are authorized.

11. Edge Cases and Business Rules

Scenario	Rule
Coverage gap	Launch is blocked. User must add/change participant context or template assignment; gap cannot be accepted.
Participant has multiple functions	AI considers all selected functions and avoids duplicate questions while preserving coverage.
Participant removed before launch	Recalculate coverage immediately.

Scenario	Rule
Participant removed after launch	Unit Admin/Project Owner must replace, cancel, or restart affected questionnaires; analysis remains blocked until complete.
Participant becomes inactive	Access is revoked; owner is prompted to replace participant; recipient delivery is blocked.
Questionnaire reopened	Prior submission is superseded, collection becomes incomplete, and downstream draft is marked stale.
Template updated after review creation	Existing instance remains on frozen version; owner may explicitly recreate/migrate an unlaunched draft.
Schedule edited	Apply prospectively to unlaunched instances; never rewrite completed history.
End date reached	No further recurring instances are generated.
AI service failure	Preserve responses; retry safely; expose status; do not silently publish partial output.
Conflicting answers	Represent conflict and evidence neutrally; do not force false consensus.
Insufficient evidence	Lower confidence and state limitation; do not fabricate a score rationale.
Approval rejected	Return to Revision Required with comments; distribution remains disabled.
Recipient outside department	Allow only if the internal user is explicitly authorized under company policy; never let recipient selection bypass project/department access rules.
Final correction	Create a new controlled report revision; retain prior final and distribution history.
Duplicate distribution	Use idempotency and log retries to avoid repeated email delivery.

12. Non-Functional Considerations

Category	Requirement direction
Security	SSO-ready identity, least privilege, encryption in transit/at rest, department-bound authorization checks, secure document scanning.
Privacy	Private participant responses, logged evidence access, configurable retention, and no external recipients in MVP.
Reliability	Durable autosave, idempotent scheduling/distribution, recoverable AI jobs, and no response loss.
Performance	Interactive pages should feel responsive; long AI/document jobs are asynchronous with progress and retry state.
Auditability	Immutable versions, actor/time/action logs, AI configuration provenance, approval record, and delivery log.

Category	Requirement direction
Accessibility	WCAG 2.1 AA target for core web flows; accessible generated PDF structure where supported.
Scalability	Partition by department/project and queue questionnaire, analysis, export, notification, and recurrence jobs.
Observability	Track job failures, latency, coverage blocks, AI exceptions, approval aging, and delivery outcomes.
Data portability	Export final PDF and retain structured report/review data for future controlled integrations.

13. MVP and Out of Scope

13.1 MVP

- Organization hierarchy, roles, department isolation, accounts/projects, and Project Owner assignment.
- Versioned questionnaire and report templates, including PDF/DOCX-to-template draft generation.
- Ad-hoc and recurring project reviews with required recurrence end date.
- Flexible participants/functions, optional AI assignment preview, and mandatory 100% coverage gate.
- Private adaptive questionnaires with bounded follow-ups and collection controls.
- Cross-participant analysis, explainable health scoring, and evidence-linked report draft.
- Human editing, mandatory Unit Admin approval, immutable final revision, internal in-app/PDF/email distribution.
- Project-level dashboards, history, notifications, and audit logging.

13.2 Post-MVP / Out of Scope

- Portfolio/account/department/company aggregate reports or benchmarks.
- Cross-period trend analysis beyond basic project history.
- External participants or recipients.
- Individual employee performance scoring.
- Advanced integrations with delivery, HR, finance, or BI systems.
- Mobile-native applications, multilingual generation, and customer-facing portals.

14. Acceptance and Success Criteria

14.1 Product acceptance criteria

Area	Acceptance criterion
Isolation	A Unit Admin or User cannot retrieve another department's protected objects through UI or API; tests cover direct-object references.
Coverage	A review with 99% or less mandatory coverage cannot launch; at 100% it can proceed if all other validation passes.
Adaptive interview	No template question receives more than 3 AI follow-ups and no participant questionnaire receives more than 10 total.

Area	Acceptance criterion
Privacy	Participants cannot view peer responses; Project Owner sees completion state during collection but not raw answers.
Versioning	Published template versions remain immutable and completed reviews render against their frozen versions.
Recurrence	A recurring schedule cannot be saved without an end date and generates no instance after that date.
Analysis	Official cross-participant analysis begins only after all required questionnaires are complete.
Approval	Finalization and all distribution controls remain disabled until Unit Admin approval is recorded.
Recipients	External/unrecognized addresses cannot be selected or emailed.
Outputs	Approved final report is available in-app and exports to PDF; email delivery and failure status are logged.
Scope	All report navigation and exports resolve to one project; no aggregate reporting is exposed.
Audit	Project edits, reopen actions, report edits, approval, finalization, and distribution are attributable and timestamped.

14.2 Initial success measures

- ≥90% of launched reviews reach 100% coverage without support intervention after normal setup.
- ≥80% of required questionnaires are completed by due date after reminders.
- Median Project Owner time from analysis completion to approval submission decreases versus the baseline process.
- ≥95% of generated reports are approved after no more than one return-for-revision cycle.
- Zero confirmed cross-department data exposures and zero final reports distributed without Unit Admin approval.
- AI follow-up cap violations: zero; external recipient deliveries: zero.

14.3 Implementation decisions to configure

- Health-score dimensions, weights, bands, minimum evidence, and confidence treatment.
- Reminder cadence, due-date defaults, escalation timing, and cancellation policy.
- Retention periods, evidence-access policy, Admin override policy, and approved internal cross-department recipient policy.
- Performance SLOs and supported file-size/page-count limits for PDF/DOCX ingestion.

15. Release Readiness Checklist

- Role and authorization tests pass across every base role and Project Owner assignment.
- Coverage engine passes mandatory-gap, multi-role, replacement, and template-version scenarios.
- AI evaluation verifies relevance, cap enforcement, privacy, evidence traceability, neutrality, and hallucination handling.
- Approval and distribution cannot be bypassed through UI, API, retries, or scheduled jobs.
- Generated reports pass content, PDF layout, accessibility, and audit-history checks.

- Operational dashboards, failure alerts, data retention, backup, recovery, and support playbooks are ready.